

Byzantine chain replication (BCR)

Pseudocode

Byzantine- Chain Replication (BCR) is a generic approach to make applications tolerate arbitrary faults beyond crash failures in an asynchronous environment. BCR algorithm tolerates a maximum of “ t ” failures of worst kind called byzantine failures. It constitutes a client performing operations at random, Olympus is the configuration service and chain of replicas which perform same set of operations and maintain similar running states.

Pseudocode defined below has a Client requesting operations to be executed at the chain of replicas and the configuration of replicas is real time managed by Olympus.

Symbols Used:

p_i : replica at i^{th} position in the configuration

C_i : current configuration (chain of replicas)

S : Slot number

o : operation

δ : cryptographic hash

Shuttle: tuple of order_proof (s,o) p_i and result_proof ($\text{result},o,\delta(r)$) p

1. Client: *#operations that take place at client side*

1. request_sequence = getRequestSequenceFromConfig() *#returns a sequence of (operation o ,requestId)*
 2. set client.timer = **Timer.start()**_{requestId} *# Timer corresponding to request Id is initiated, maintains bound on turn around time.*
 3. **client** sends configuration request to **Olympus**
(replica_sequence, replica_keys) = getCurrentConfigFromOlympus(); *# returns a current replica configuration if there exists one else returns a new configuration.*
-

#After execution of section 1 involving client operations, olympus continued below.

2. Olympus: On receiving `getCurrentConfigFromOlympus()`

1. **If**(configuration does not exists):
 - a. Generate a configuration C_i with p replicas
 - b. **for** all replica **in** Configuration C_i
inithist(); *#inithist is the valid history of operation per slot | running state of replica. Empty initially.*
`replica_key = generate_Key(replica):`
#replica_key.public = public_key of replica
#replica_key.private = private_key of replica
 - c. Broadcast corresponding public keys to all the replicas.
2. **else return** (C_i (configuration)) to the client.

#After execution of section 2 involving Olympus, Client operations continued below.

3. Client : On receiving (C_i (configuration)):

1. From the $C_i(p_1, p_2 \dots p_n)$, Fetch the **head_replica**. *# where C_i is the current configuration*
`received_result (order_proof, result_proof) = send_operations_to_replica(operation, *head)`
#for send_operations_to_replica goto replica section 4.
*# *head is the pointer to the head replica*

Where:

order_proof is signed order statement $(s, o)p_i$ for all p_i belongs to C .

result_proof is signed tuple of $(result, o, \delta(r))p_i$. $\delta(r)$, $\delta(r)$ is encrypted replica result.

Case: fault-free

- i. **if** (`hash(result(r)) == $\delta(r)$`) *# hash is applied by client on result to verify if $\delta(result)$ from replica matches*
“successfully executed operation”
else *# (`hash(result(r)) != $\delta(r)$`) hash result did not match the result returned by replica, this constitutes as proof of misbehaviour, hence reconfiguration request sent to Olympus*
reconfigure_replicas(proof_of_misbehaviour) *#declared in the end*

Case: failure.

if (client.timer == expired): *# timer was initiated in step 2.*

a. received_result = **broadcast_operations**(operation(o), replica_Sequence)

*#broadcast_operations retransmits request to all the replicas. Function declared below
under replica(faulty case):*

*#received_result captures the result of all replicas, once client **retransmit** operation to all
replicas in current config.*

i. if (received_result == **ERROR**): *# this means that the replica is immutable*

- Client requests latest configuration from Olympus.

replica_sequence = **getCurrentConfigFromOlympus**();

*# returns a current replica configuration if there exists one else
returns a new configuration.*

- Client executes request protocol from scratch. *#from section 3*

ii. elseif (received_result == result(r))

if (hash(result(r)) == delta(r)) *# replica has the result cached, client
receives it and verifies*

“successfully executed operation”

else if (hash(result(r)) != delta(r)): *# hash result did not match the
result returned by replica, this constitutes as **proof of misbehaviour**, hence reconfiguration
request sent to Olympus*

reconfigure_replicas(proof_of_misbehaviour) *#declared in end*

iii. elseif(received_result = Error and result(r)) *# some replica returned **Error**,
some returned **result(r)***

if (hash(result(r)) == delta(r)):

“successfully executed operation”

else if (hash(result(r)) != delta(r)):

reconfigure_replicas(proof_of_misbehaviour)

4. Replica : *#on receiving send_operations_to_replica(operation, *head) from client*

1. p_{head} orders the operations and assign slot(s) number.
2. **for** p_i in (Ci(configuration)): *# iterating from head to tail request protocol goes through every replica in configuration, executes the operation and updates the order_proof, result_proof*

if($p_i == head$):

- a. create **shuttle** containing ((s,o) p_i , (result,o,delta(r)) p_i) *# head replica creates a shuttle containing order_proof (s,o) p_i and result_proof (result,o,delta(r)) p_i .*

- b. Increment **checkpointing_counter**: *#counter to implement checkpointing*

if(**checkpointing_counter** mod N == 0): *#when checkpointing*

counter is a multiple of N checkpoint shuttle is along the chain.

Checkpointing(): *# the function handles checkpoint algorithm and is declared towards the end.*

else:

for all p_i in replica_sequence: *#starting from head to tail*

- i. **for all** $p_j < p_i$ in replica_sequence:

- a. **Validate order_proof** (<order,s,o>) p_j *#for all predecessor replica p_j , (s,o) p_j should belong to (<order,s,o>) p_i*

- b. **execute** operation o and obtain **result** r.

- ii. add its (s,o) p_i to the p_i .history which is a list of **order_proof**.

- iii. add its result (result,o,delta(r)) p_i to its **result_proof**.

- iv. **if** ($p_i == tail$): *# if shuttle reaches the tail replica*

return **result_proof** (result,o,delta(r)) p_i to client.

3. **for** p_i in (Ci(configuration)): *#iterating from tail to head replica.*

if ($p_i == tail$):

- a. Tail forward the **result shuttle** ((s,o) p_i , (result,o,delta(r)) p_i) back to predecessor replicas.

- b. **if**(p_i .timer_requestId is active): *#timer against requestId was started*

this is the case where timeout occurs at client, operation is broadcasted to replicas, if operation is not recognized protocol is executed from scratch. Cancel the timer in case result_shuttle is received before timeout. Otherwise send a reconfiguration request to olympus.

if p_i receives the **result shuttle**: *#in case of retransmission and timer against the received request ID is active*

- **cancel the timer**
- **cache** $result_proof = (result, o, \delta(r))p_i$
- **return** $result_proof$ *#to the client*

for any p_i **in** $replica_sequence$:

- **if** (p_i timer == expire):
 reconfigure_replicas() *#send reconfiguration request to olympus. Function declared towards the end.*

c. **cache**(key=requestId, value=result_shuttle) *# cache the result replica with key as requestId and value as the result_replica*

else: *# replicas other than tail*

cache(key=requestId, value=result_shuttle) *#cache the result*

4. Replica : in case of timeout failure : **broadcast_operations** received from **client**

case when timeout happens at client

broadcast_operations called from section 3

broadcast_operations(operation(o), replica_sequence): *# where o is operation*

for all p_i **in** $replica_sequence$:

if ($p_i.mode = \text{Immutable}$)

return “Error” statement.

if ($p_i.mode = \text{Active}$)

if (**cached_result**(requestId) **exists**) *# cached result is present against the requestId*

return $result_proof = (result, o, \delta(r))p_i$ *# returns cached result*

else *# if no cached result found*

- $p_i.timer = \text{Timer.start}()_{o.requestId}$ *# starting timer corresponding to requestid of o at replica meanwhile request forwarded to head*
- pass the request to p_{head} .
- p_{head} upon receiving the request:

 1. **if** (**cached_result**(requestId) == exists) :

return $(result, o, \delta(r))p_{head}$ *# this is the result_proof*

 2. **elseif** (**operation**(o) is processing and **result_shuttle** yet to receive))

```

Phead.timer = Timer.start()requestId #starting timer at head
replica corresponding to the requestId of the operation.
if( result shuttle is received before timer Phead.timer expires ):
    i. pi.timer = Timer.cancel() #cancel timer at replica
    ii. cache the result_proof
    iii. return (result,o,delta(r))phead # this is the
        result_proof
else: (Phead.timer == expired):
    reconfigure_replicas(proof_of_misbehaviour)
    
```

```

3. elseif (operation o is not recognized):
    send_operations_to_replica(operation, *head()* # starts
    the protocol from scratch, goto section 4 Replica
    
```

4. Replica:

Case : when order proof not validated | peer faulty | timeout at client

1. **If** (**p_i** finds peer faulty || **order_proof** not validated)
 - i. Send wedge statement(**<wedge, h>p_i**) to **Olympus** *#h denotes the history which contains order proofs*
 - ii. **P_i.mode=Immutable** *# The replica changes it's state to immutable, i.e. it can't issue new order statements*

Olympus side:

On receiving wedge statement(**<wedge, h>p_i**)

Broadcasts **wedge** request to all replicas *#broadcast wedge request to all replicas requesting their history*

Replica side:

On Receiving wedge request from Olympus

for all replica **p** in (Ci(configuration)):

Send wedge statement and order proof (**<wedge, h>p_i** & **<order,s,r>p_i**) to **Olympus**

Olympus Side:

On receiving ($\langle wedge, h \rangle_{p_i}$ & $\langle order, s, r \rangle_{p_i}$) *statement from replicas*

#the replicas in the current config will have at least 1 honest replica among max T failures

for all replicas check

if(combination of ($\langle wedge, h \rangle_{p_i}$) wedge statement and $\langle order, s, r \rangle_{p_i}$ statement

does not includes a order_proof for (s,o)) *# this concludes as proof of misbehaviour and reconfiguration function is called*

reconfigure_replicas(proof_of_misbehaviour) *#declared in the end*

reconfigure_replicas(): *#reconfiguration algorithm*

Olympus: On receiving ($\langle wedge, h \rangle_{p_i}$)_{C_i} from all the replicas

1. let Q be quorum of replica in C_i with valid histories
2. for every pair p_i and p_j in Q
 - a. if there exists a $\langle slot, operation1 \rangle$ in $p_i.history$ and $\langle slot, operation2 \rangle$ in $p_j.history$
if (operation1 == operation2):
Then slot history is consistent
else: Choose different Quorum and goto step-1 *#(repeat reconfiguration algorithm)*
 - b. LH = longest of the $\langle slot, operation \rangle$ pairs from all replica corresponding to a slot number *#LH is longest history*
3. for all replica(p) in Q:
 - a. (catch_up) p_i = (LH - $p_i.history$) *#suffix of LH that p_i has not executed yet*
 - b. send (catch_up) p_i to replica in Q *#olympus sends catch up message to replica*

On Replica Side:

4. for all replica p_i on receiving (catch_up) p_i message
 - a. p_i executes operations in (catch_up) p_i
 - b. $ch = \text{delta}(\text{running_state})_{p_i}$ *#ch is the cryptographic hash of p_i 's running state S*
 - c. send(caught_up(ch)) to Olympus

On Olympus side:

- a. after receiving "caught_up" from all replicas in Q
for replica in quorum Q:
if(replica "caught_up" != ch) *# checking if any replica "caught_up" value is not ch*
choose different Quorum and goto Step-1 *#(repeat reconfiguration algorithm)*

```
else #all replicas have "caught_up" value = Ch
    i. Send "get_running_State" message to any replica in Q
    # replica returns its running state S
    if(delta(S) != ch)
        Request running_State from another replica in Q
    else:
        i. Generate a configuration Cj with p replicas
        ii. for all replica in Configuration Cj
            inithist(S); # S is the current running state of the previous
            configuration
```

On receiving *Checkpointing()* from replica in section 4.

Checkpointing(): #*(at Nth operation)*

while(checkpointing_counter mod N == 0): #*function is called after every N operations*

1. for all p_i in replica_sequence: #*iterating from head replica to tail replica.*
 - a. if ($p_i == \text{head}$):
initiate a shuttle containing the checkpoint and a running state.
(checkpoint, delta(running_state)) p_i #*checkpoint_proof tuple*
 - b. else:
add a (checkpoint,delta(state)) p_i to the checkpoint_proof.
2. for all p_i in replica_sequence: #*iterating tail replica to head replica.*
 - truncate prefix from p_i history
 - add a checkpoint in history corresponding to deleted history.
 p_i history becomes ((s, o, p_i , C, (checkpoint,order_proof) p_i)
once checkpoint_shuttle is received remove the prefix from history until the checkpoint.
 - return the shuttle to the next replica in sequence towards head.

Bibliography

1. Byzantine Chain Replication research paper by *Van Renesse, Chi Ho, and Nicolas Schiper.*
2. Project.txt document provided by *Prof. Scott Stoller*